



# ADACS

ASTRONOMY DATA AND COMPUTING SERVICES

## Astroinformatics school - "Rise of the machines"



4 to 6 February 2019

Presented by Rebecca Lange and Dan Marrable

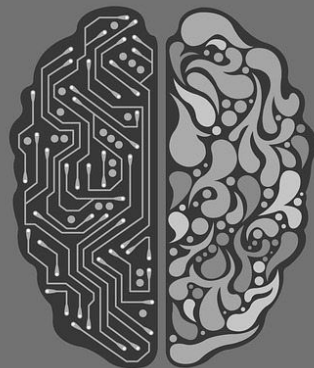
# Artificial Neural Networks



Curtin Institute for Computation

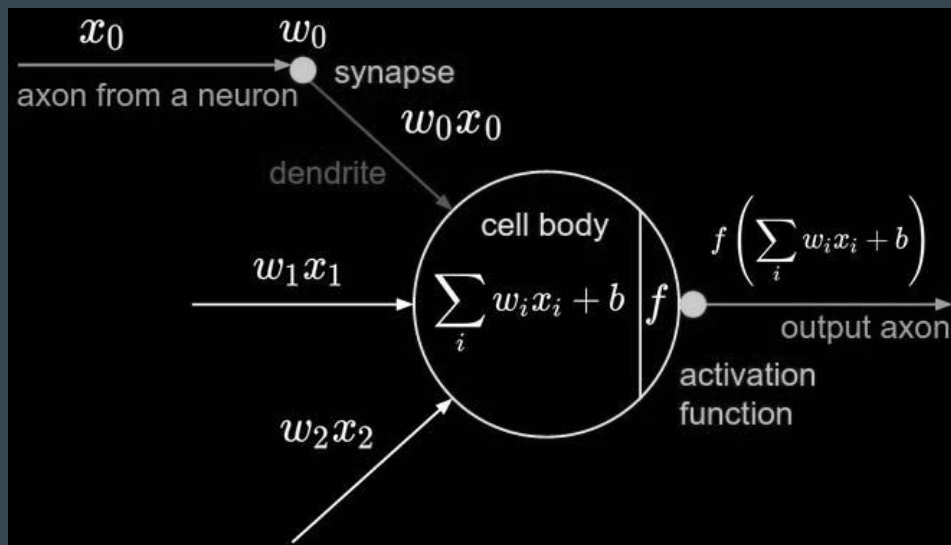
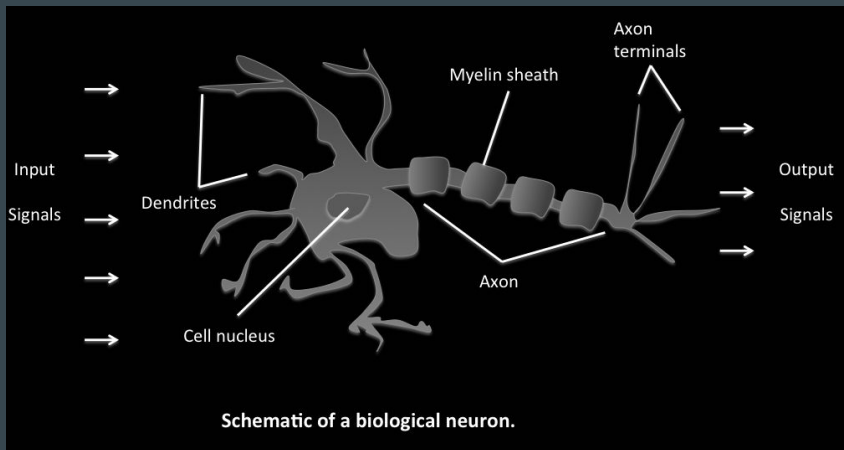
# What are Artificial Neural Networks?

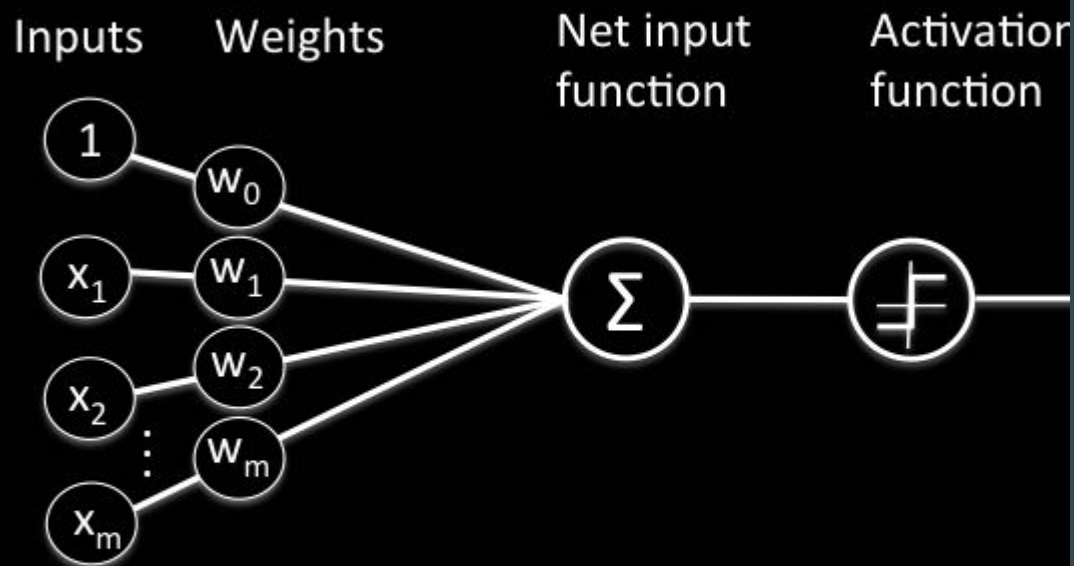
- Universal function approximators
- A series of linear equations (and activation functions) that approximate nonlinear equations.
- Modelled loosely around how the brain's synapses works.
- Generally used for classification tasks



# Neuron

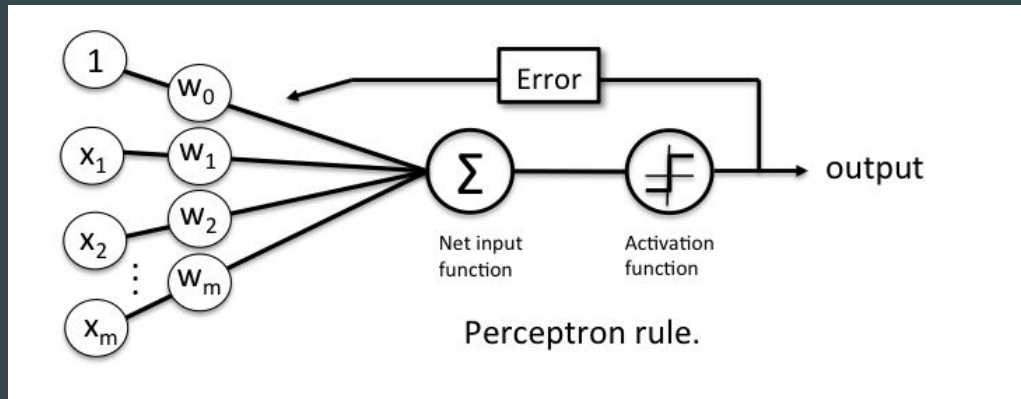
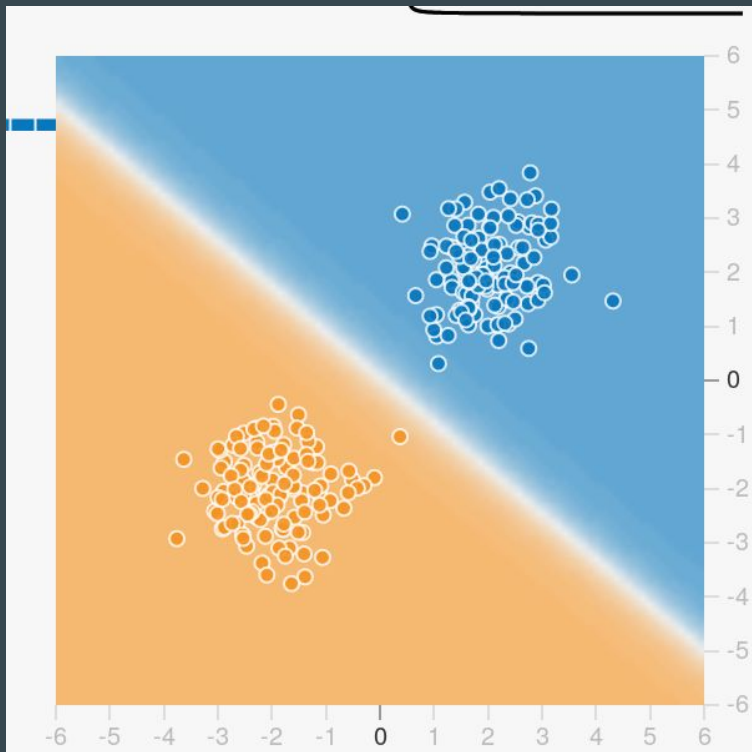
# Biological Inspiration





**Schematic of Rosenblatt's perceptron.**

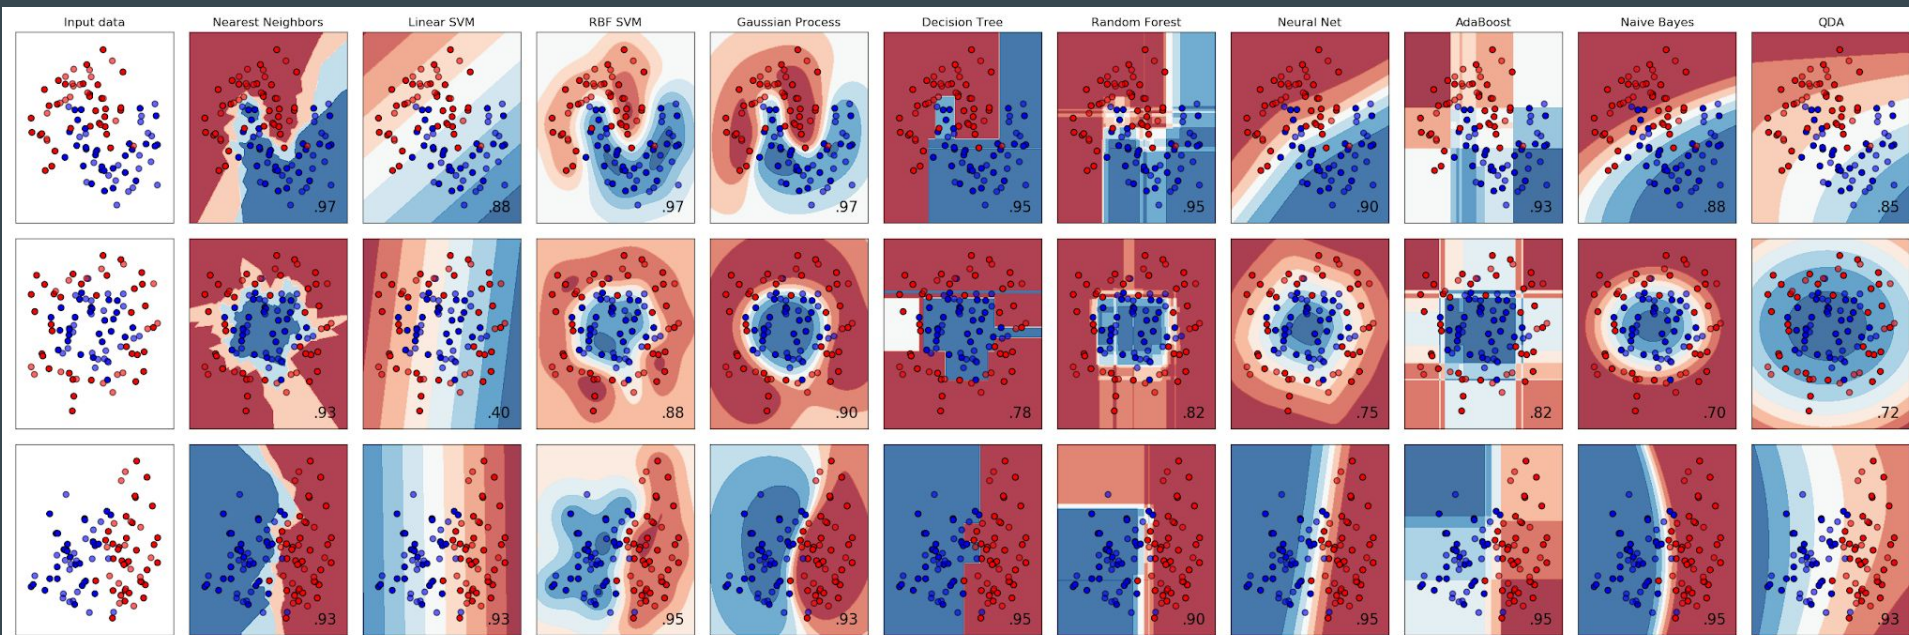
# Single Perceptron



# Limitations

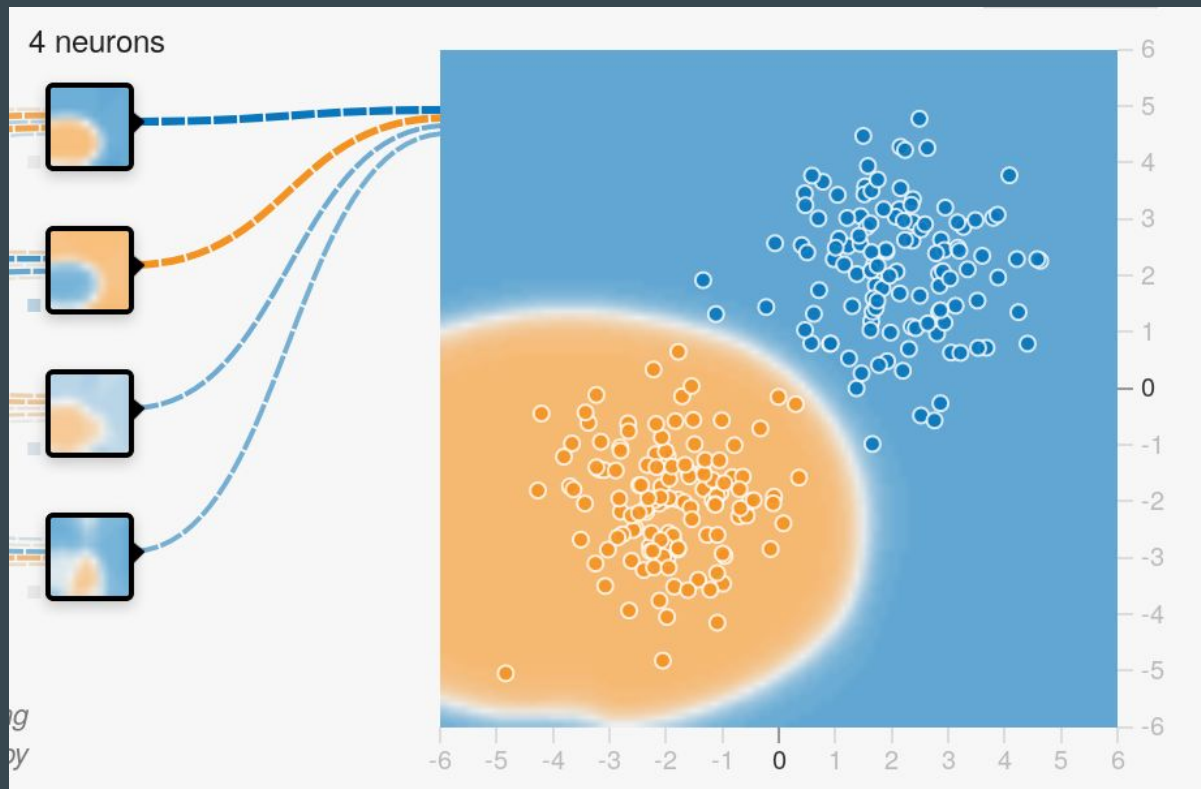
A perceptron can only be used to linearly separable classes

*None of these cases are possible*





# A Network of Neurons

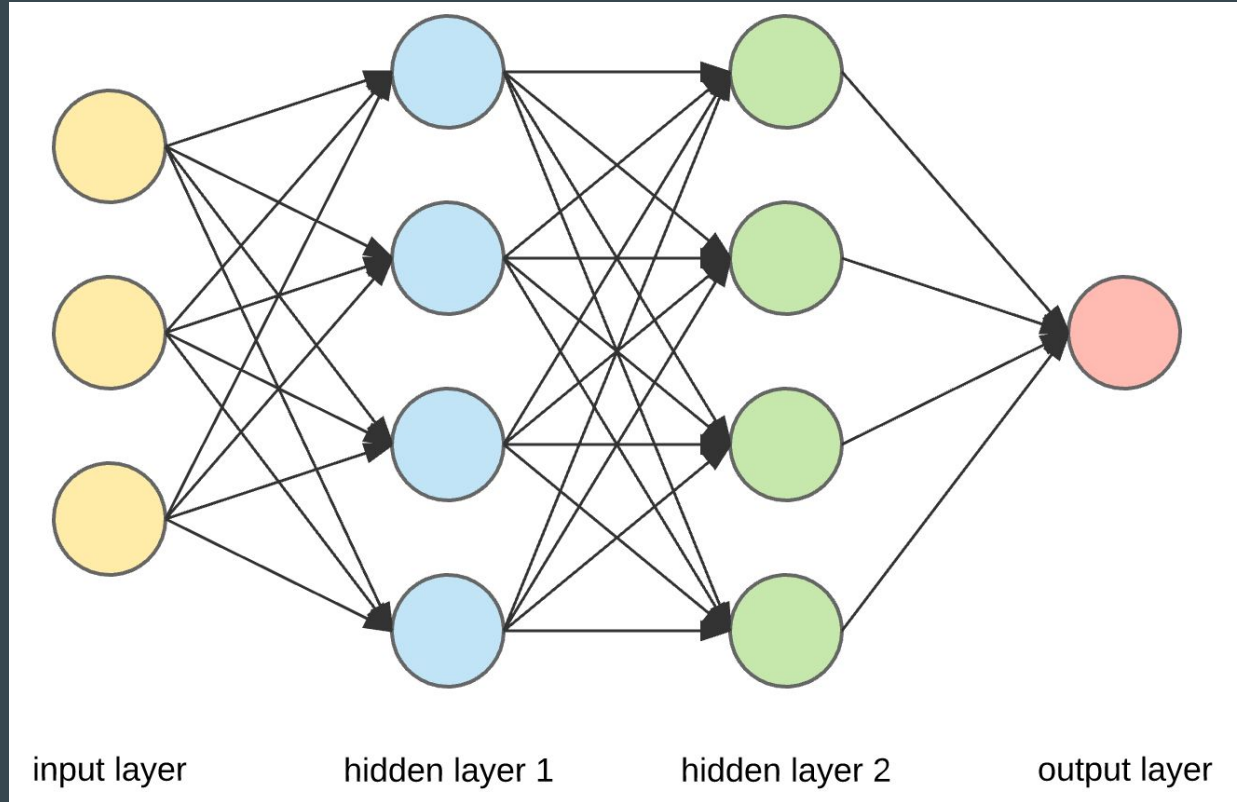


# A Typical Neural Network

Input Layers takes in the data

Hidden layers have most of the training parameters

The output layer gives the classification



# Example

Some combination of parameters will make our output layer activate correctly

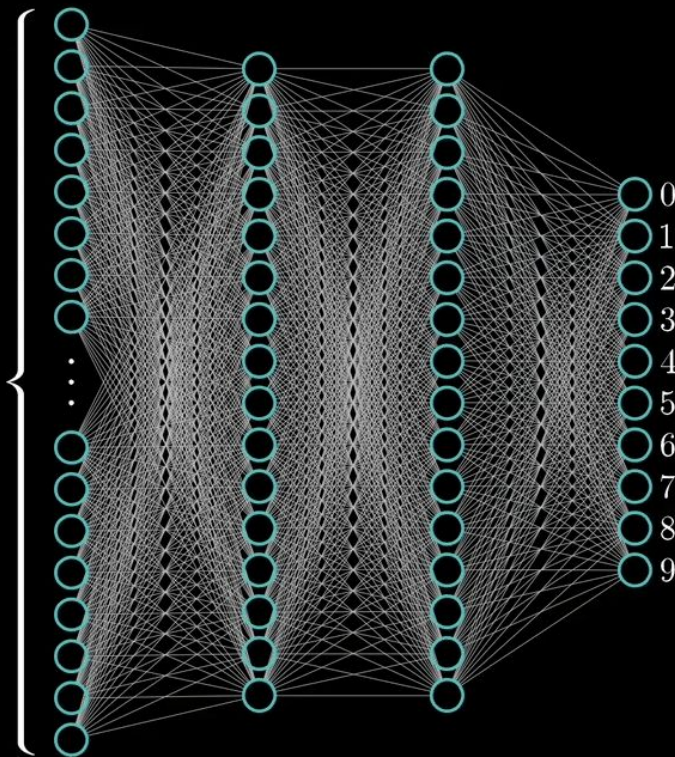
We don't know what they are yet

There are too many to randomly guess

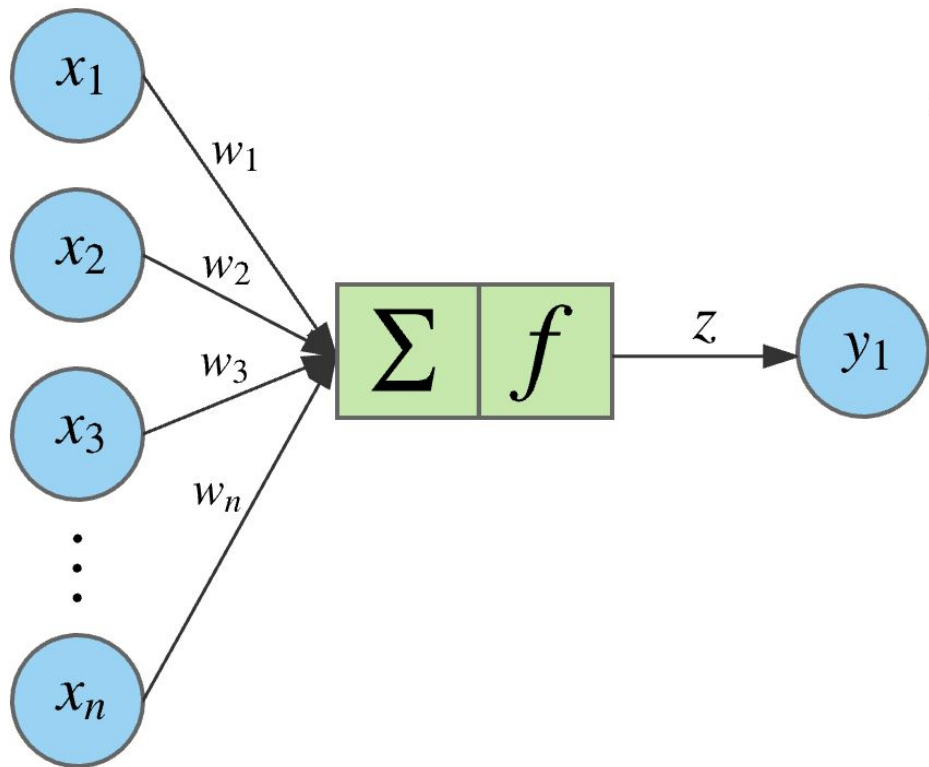
We need an algorithm to find them



784



# Mathematical Representation



$$z = f(b + x \cdot w) = f \left( b + \sum_{i=1}^n x_i w_i \right)$$

$$x \in d_{1 \times n}, w \in d_{n \times 1}, b \in d_{1 \times 1}, z \in d_{1 \times 1}$$

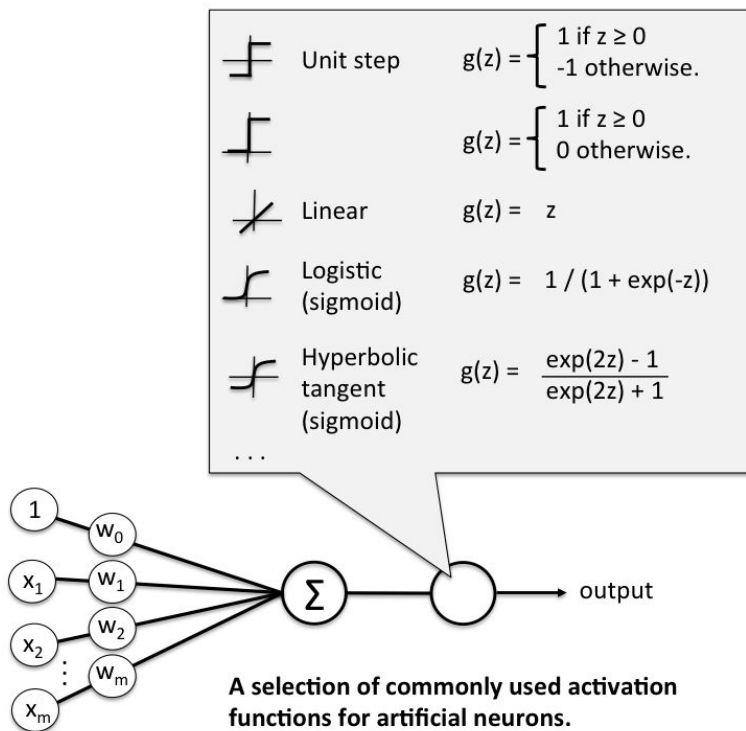
Multiply the inputs by a **weight**

Sum the result

Add a **bias**

Outputs are passed to an **activation function**

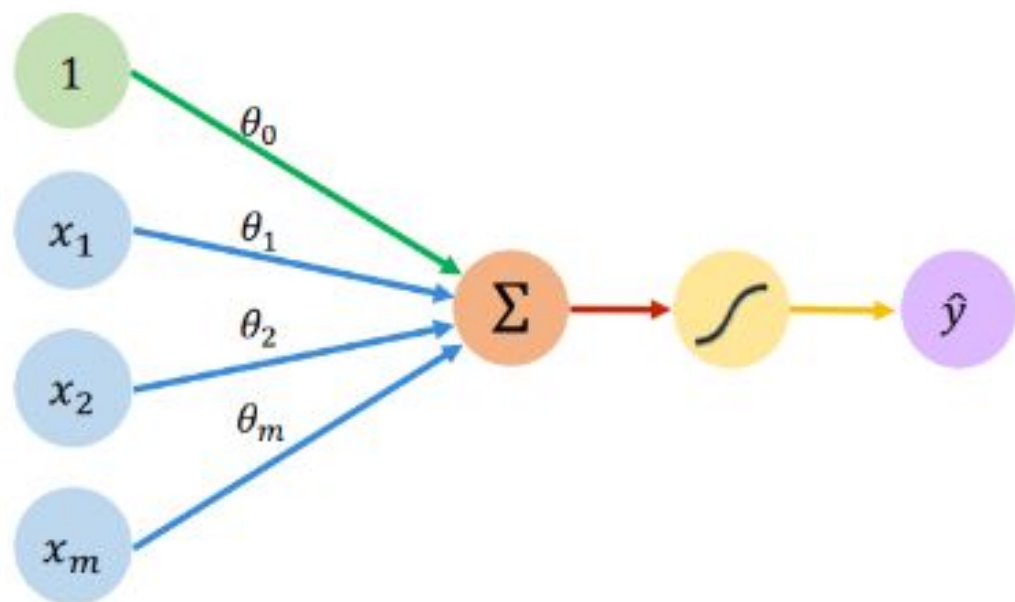
# Activation Function



Activation functions give the neurons a non-linear response

Sigmoids converge to either 1 or 0, but are slow to train and 'kill' gradients

If unsure, the rule of thumb is, use ReLU for hidden layers and sigmoid for the output layer



Inputs      Weights      Sum      Non-Linearity      Output

Output

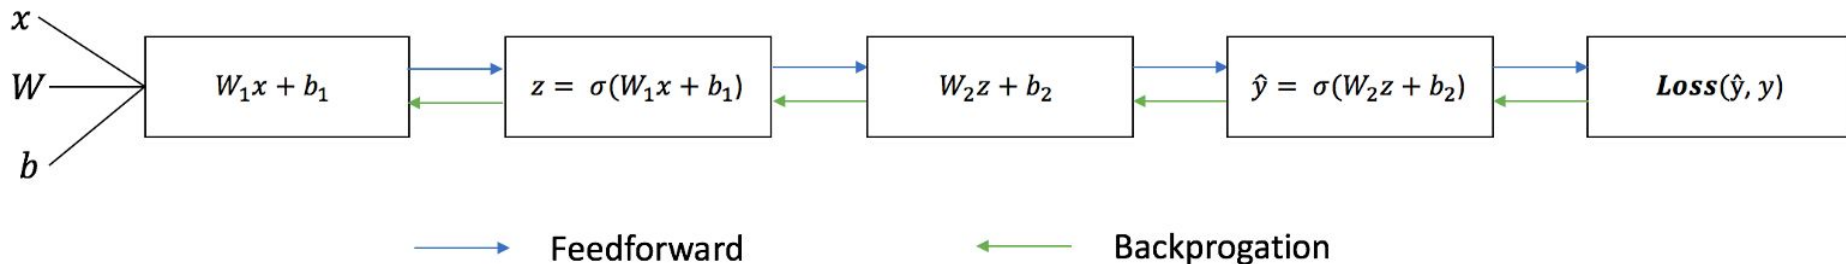
Linear combination of inputs

$$\hat{y} = g \left( \theta_0 + \sum_{i=1}^m x_i \theta_i \right)$$

Non-linear activation function

Bias

# Training our ANN



There are different loss functions.

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multilabel classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse
Regression to values between 0 and 1	sigmoid	mse or binary_crossentropy

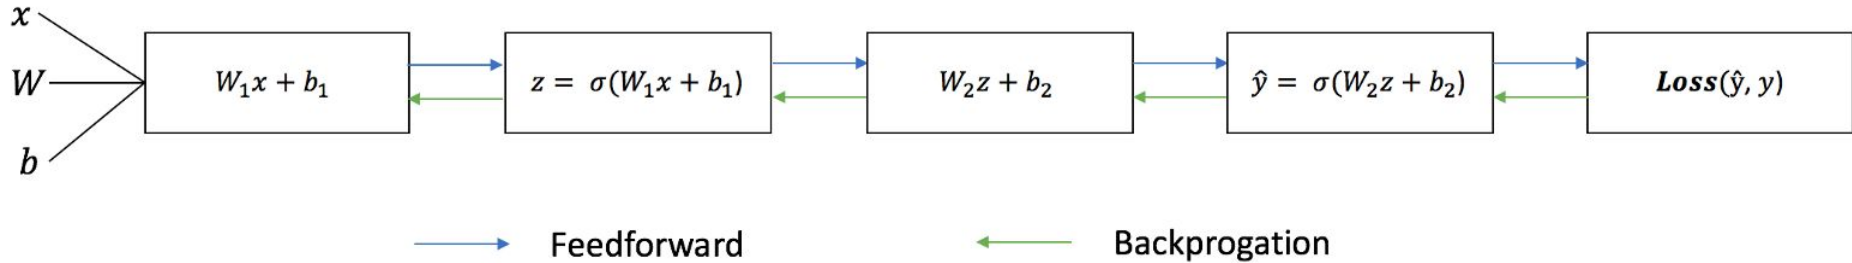
Make a prediction by making a forward propagation

Calculate the Loss

Back propagation to update the training parameters

# Training our ANN

Full vectorised derivation  
(<https://www.coursera.org/specializations/deep-learning>)



$$Loss(y, \hat{y}) = \sum_{i=1}^n (y - \hat{y})^2$$

$$\frac{\partial Loss(y, \hat{y})}{\partial W} = \frac{\partial Loss(y, \hat{y})}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z} * \frac{\partial z}{\partial W} \quad \text{where } z = Wx + b$$

$$= 2(y - \hat{y}) * \text{derivative of sigmoid function} * x$$

$$= 2(y - \hat{y}) * z(1-z) * x$$

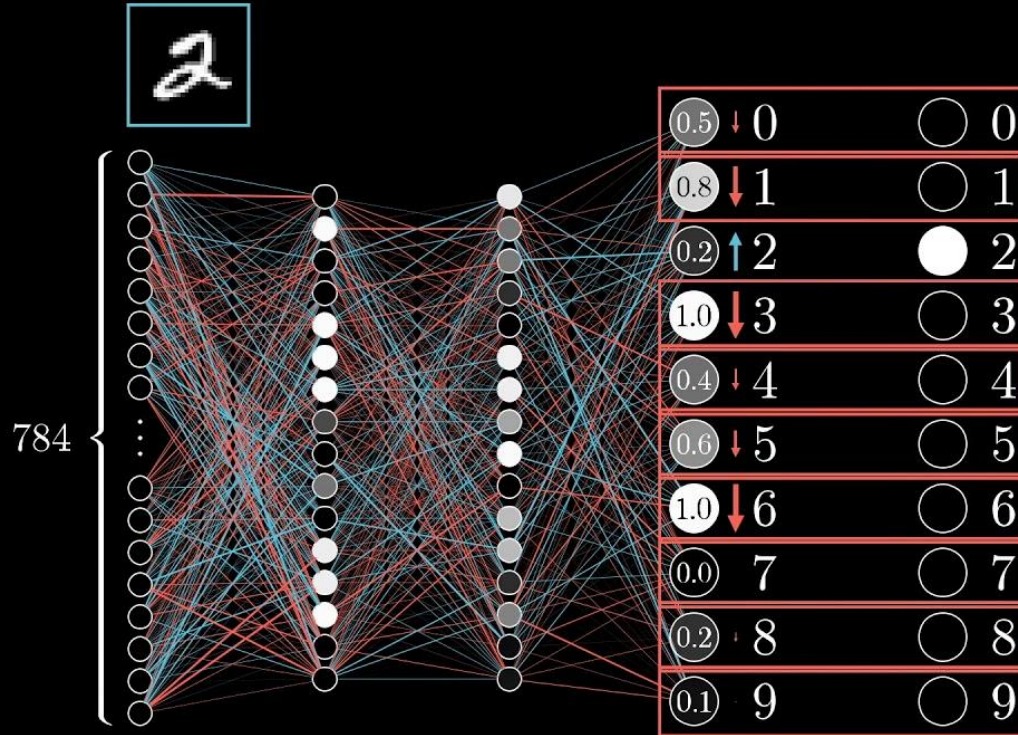
$$\Delta W_i = \eta \frac{\partial E}{\partial w_i}$$

$\eta$  is the learning rate

$$W_i = W_i + \Delta W_i$$



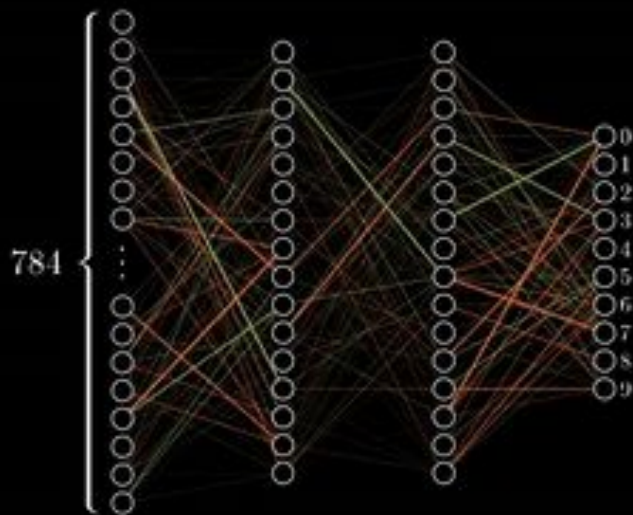
# Backpropagation



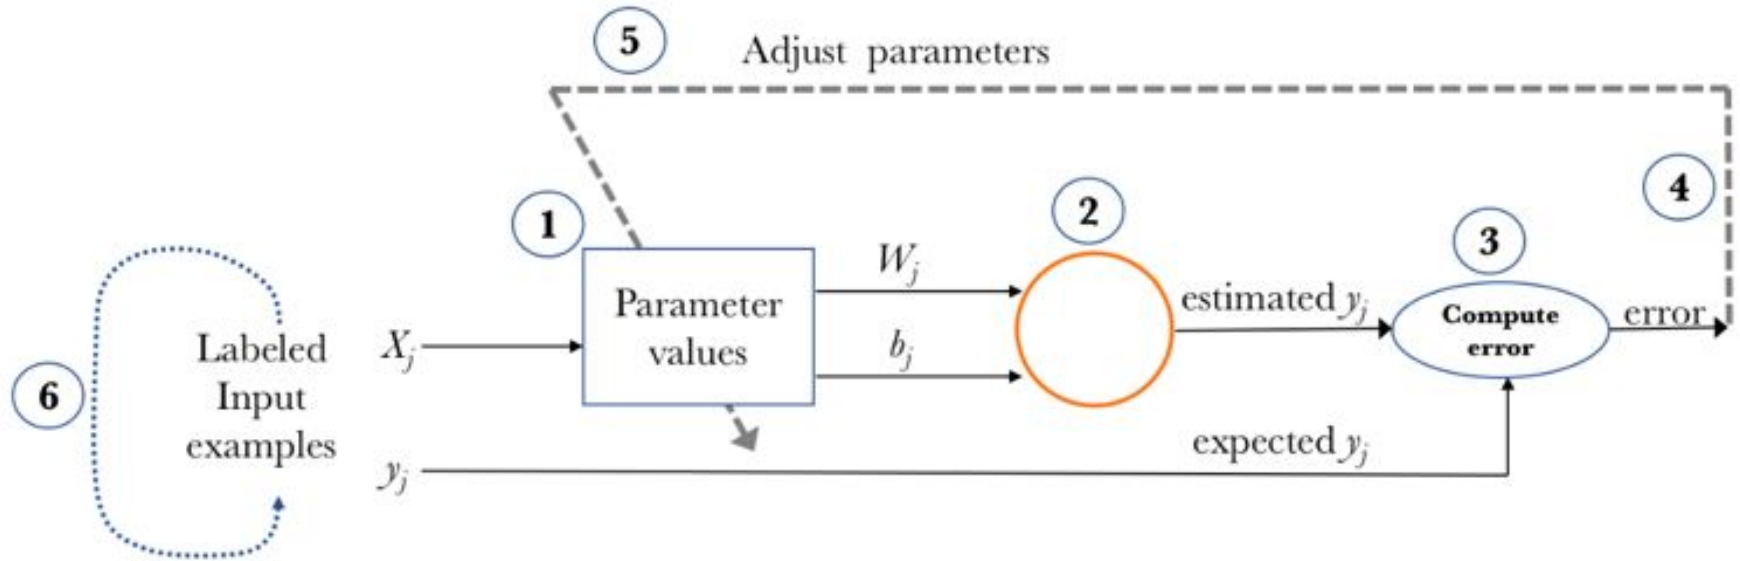
# Training the model

3Blue1Brown : <https://youtu.be/aircAruvnKk>

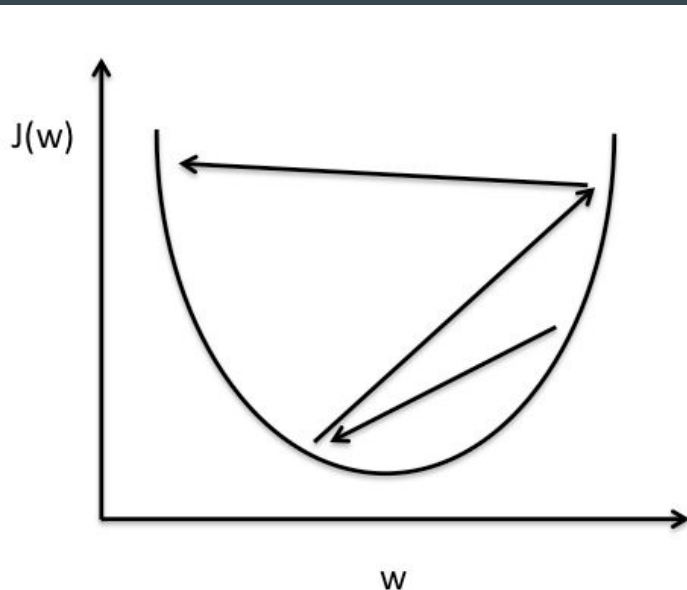
Training in  
progress...



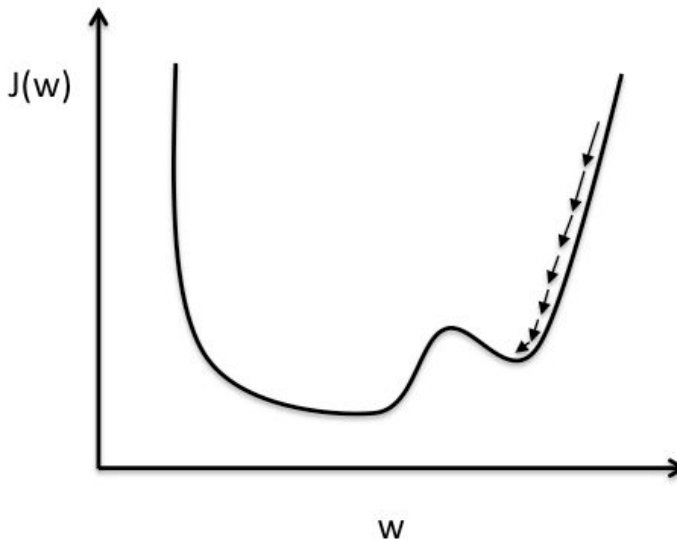
# Optimiser - predictor - corrector



# Choosing the right learning rate



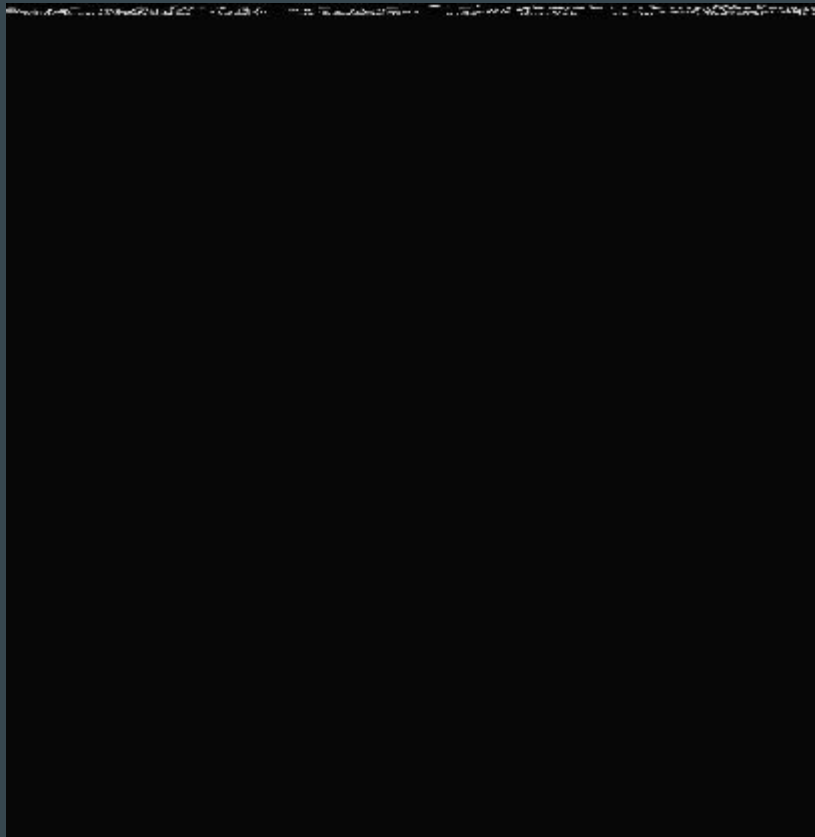
**Large learning rate: Overshooting.**



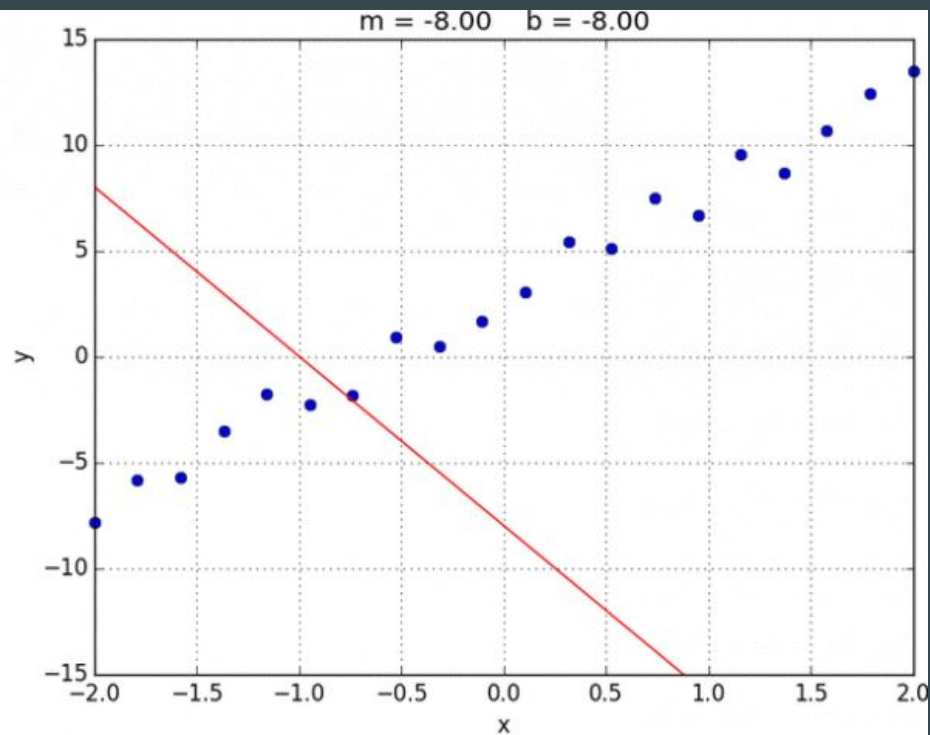
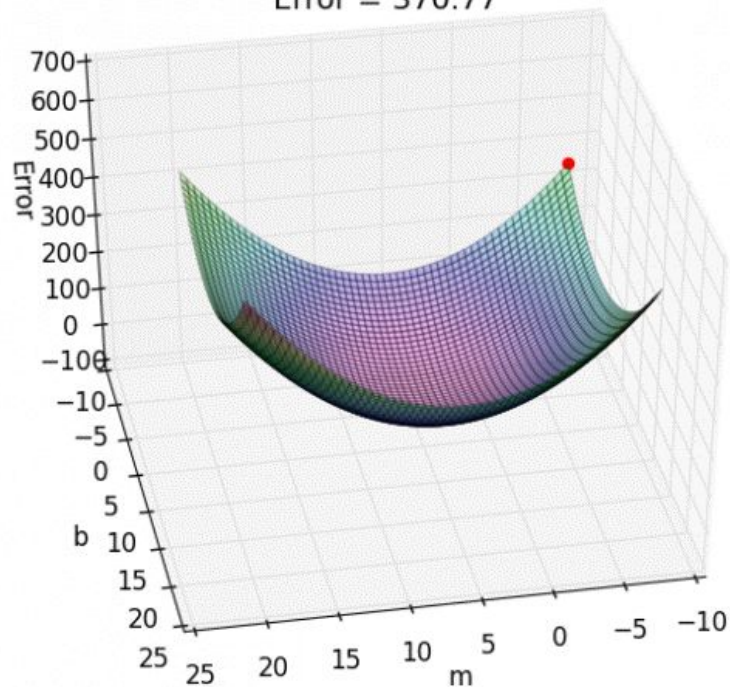
**Small learning rate: Many iterations until convergence and trapping in local minima.**

# Gradient descent

- Randomize our weights
- Perturb our weights
- Calculate the gradient of our loss
- Update the weights
- Repeat for a number of epochs



Error = 370.77



# Recap

Activation function ?

Loss function ?

Optimizer ?

Gradient descent ?

Epochs ?

Batch size ?

Learning Rate ?

Hyperparameters ?

# Recap

Activation function - Makes our layer non-linear

Loss function - Tells us how far off our prediction was

Optimizer - Algorithm that defines the learning process

Gradient descent - One method used to update the training parameters

Epochs - The number of iterations we use to optimise the model

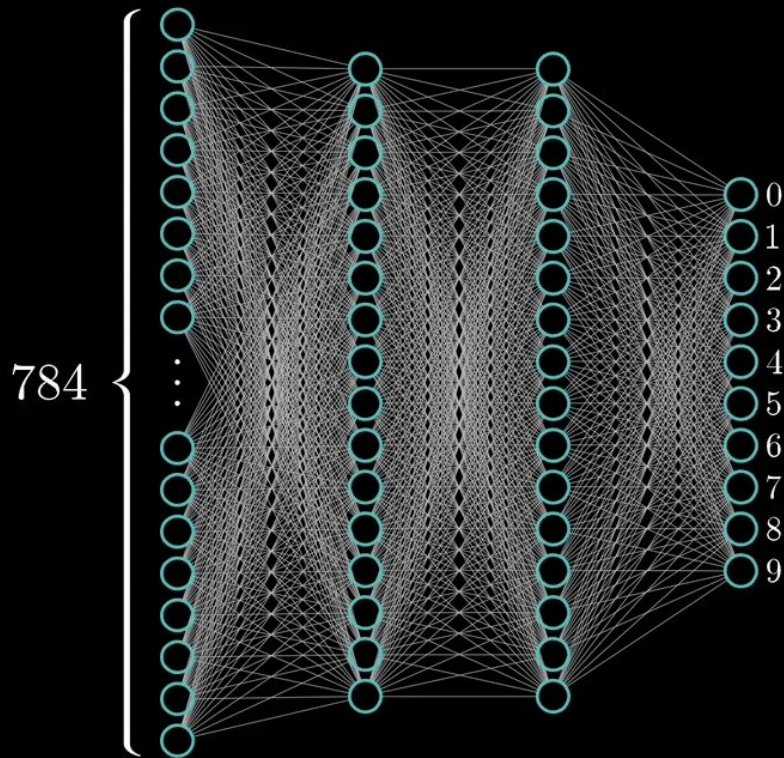
Batch size - The number of forward passes made on multiple data before updating the training parameters

Learning Rate - How great a step taken to update the training parameters

Hyperparameters - Non-trainable parameters



# Questions ?



# Web Links

- <https://towardsdatascience.com/applied-deep-learning-part-1-artificial-neural-networks-d7834f67a4f6>
- [http://sebastianraschka.com/Articles/2015\\_singlelayer\\_neurons.html](http://sebastianraschka.com/Articles/2015_singlelayer_neurons.html)
- <https://www.youtube.com/watch?v=aircAruvnKk>
- <https://www.coursera.org/specializations/deep-learning>
- <https://alykhantejani.github.io/images/>