



# ADACS

ASTRONOMY DATA AND COMPUTING SERVICES

## Astroinformatics school - "Rise of the machines"



4 to 6 February 2019

Presented by Rebecca Lange and Dan Marrable

# Convolutional Neural Networks



Curtin Institute for Computation

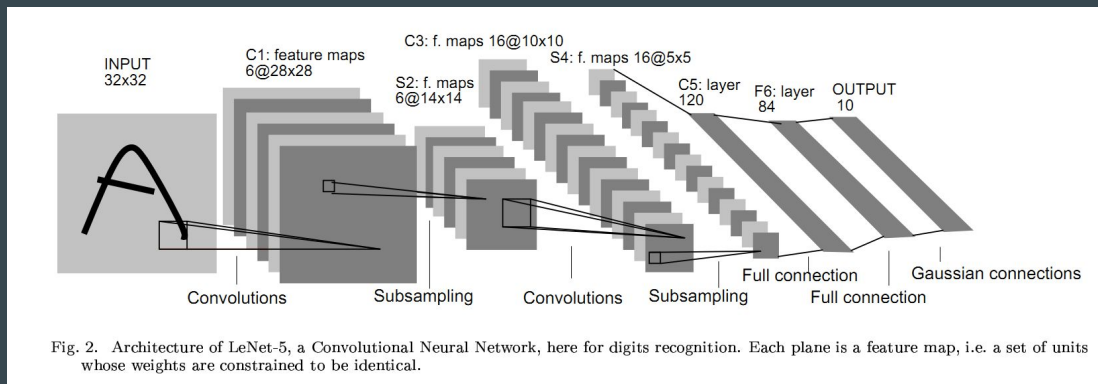
# What is a convolutional neural network (CNN)

Most commonly used for  
analyzing visual imagery

Image classification

Object detection


Image segmentation



LeNET-5 One of the first ever CNNs


# What is a convolution?

A convolution is an integral that expresses the amount of overlap of one function as it is shifted over another function . It therefore "blends" one function with another.



$$\begin{aligned}(f * g)(t) &\stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau \\ &= \int_{-\infty}^{\infty} f(t - \tau)g(\tau) d\tau.\end{aligned}$$

# Discretised Convolution in 2D



$f_1$	$f_2$	$f_3$
$f_4$	$f_5$	$f_6$
$f_7$	$f_8$	$f_9$

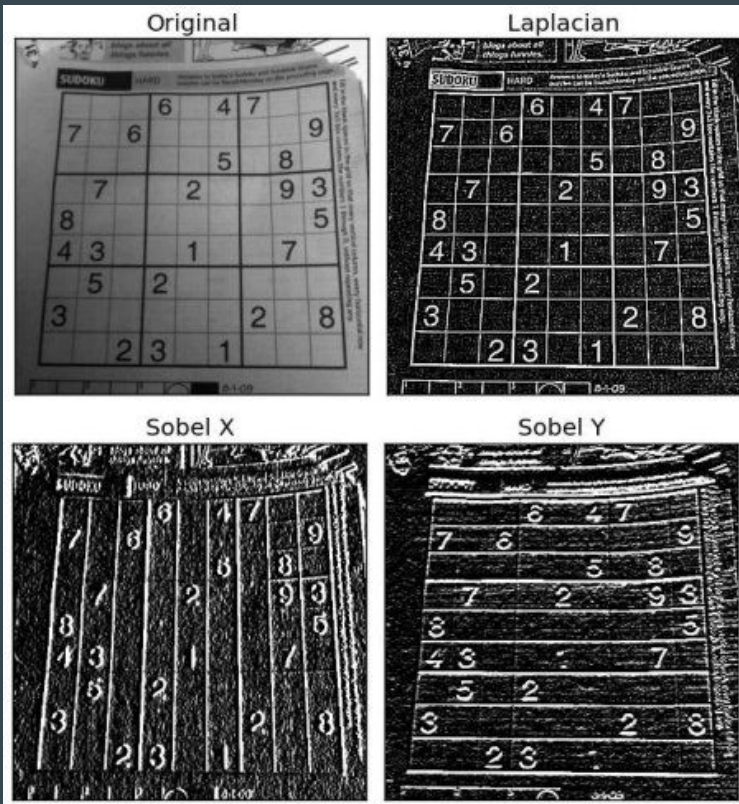
$\otimes$

$h_9$	$h_8$	$h_7$
$h_6$	$h_5$	$h_4$
$h_3$	$h_2$	$h_1$

$$\begin{aligned}
 f * h = & f_1 h_9 + f_2 h_8 + f_3 h_7 \\
 & + f_4 h_6 + f_5 h_5 + f_6 h_4 \\
 & + f_7 h_3 + f_8 h_2 + f_9 h_1
 \end{aligned}$$

# Image filters

Features



Filters

Laplacian

0	-1	0
-1	4	-1
0	-1	0

Sobel X


-1	0	+1
-2	0	+2
-1	0	+1

Gx

Sobel Y

+1	+2	+1
0	0	0
-1	-2	-1

Gy



Input Volume (+pad 1) (7x7x3)

 $x[:, :, 0]$ 

0	0	0	0	0	0	0
0	0	0	1	0	2	0
0	1	0	2	0	1	0

0	1	0	2	2	0	0
0	2	0	0	2	0	0
0	2	1	2	2	0	0
0	0	0	0	0	0	0

 $x[:, :, 1]$ 

0	0	0	0	0	0	0
0	2	1	2	1	1	0
0	2	1	2	0	1	0

0	0	2	1	0	1	0
0	1	2	2	2	2	0
0	0	1	2	0	1	0
0	0	0	0	0	0	0

 $x[:, :, 2]$ 

0	0	0	0	0	0	0
0	2	1	1	2	0	0
0	1	0	0	1	0	0

0	0	1	0	0	0	0
0	1	0	2	1	0	0
0	2	2	1	1	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

 $w0[:, :, 0]$ 

-1	0	1
0	0	1
1	-1	1

 $w0[:, :, 1]$ 

-1	0	1
1	-1	1
0	1	0

 $w0[:, :, 2]$ 

-1	1	1
1	1	0
0	-1	0

Bias b0 (1x1x1)

 $b0[:, :, 0]$ 

1
---

Filter W1 (3x3x3)

 $w1[:, :, 0]$ 

0	1	-1
0	-1	0
0	-1	1

 $w1[:, :, 1]$ 

-1	0	0
1	-1	0
1	-1	0

 $w1[:, :, 2]$ 

-1	1	-1
0	-1	-1
1	0	0

Bias b1 (1x1x1)

 $b1[:, :, 0]$ 

0
---

Output Volume (3x3x2)

 $o[:, :, 0]$ 

2	3	3
3	7	3
8	10	-3


 $o[:, :, 1]$ 

-8	-8	-3
-3	1	0
-3	-8	-5

toggle movement




# Learning Features





Alexnet Filters



Input








Input Layer

Hidden Layer


# What Do the Layers See?

## Feature Visualization

How neural networks build up their understanding of images



**Edges** (layer conv2d0)



**Textures** (layer mixed3a)



**Patterns** (layer mixed4a)



**Parts** (layers mixed4b & mixed4c)




**Objects** (layers mixed4d & mixed4e)


Feature visualization allows us to see how GoogLeNet [1], trained on the ImageNet [2] dataset, builds up its understanding of images over many layers. Visualizations of all channels are available in the [appendix](#).




# Convolutions on the hidden layers




# Strides




Stride 1



Feature Map

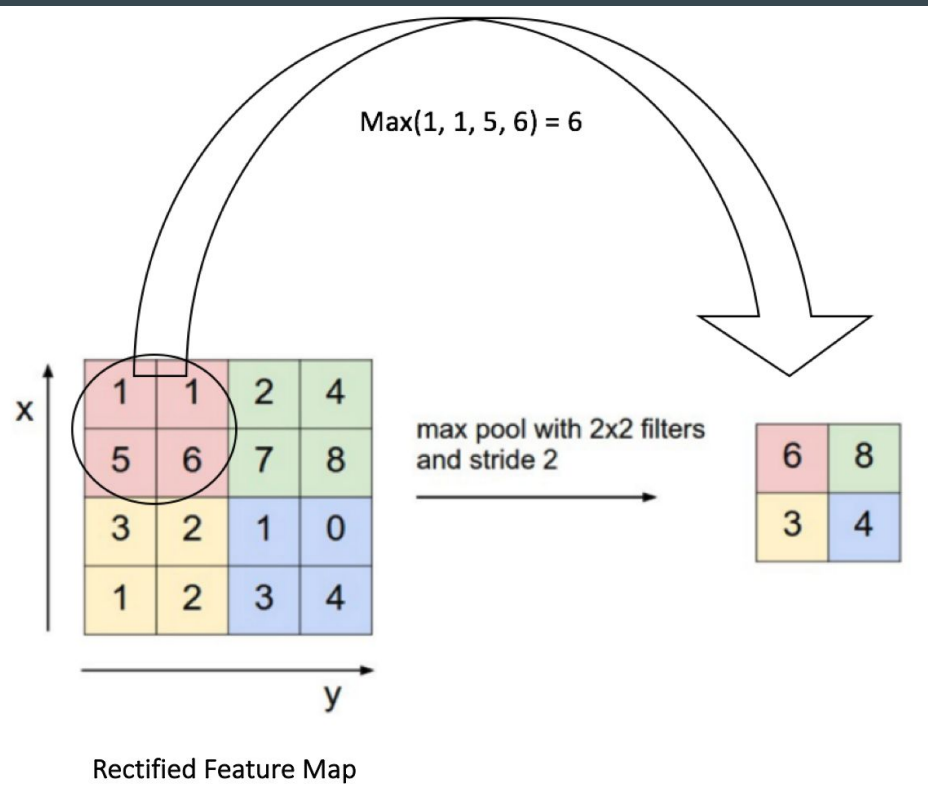


Stride 2



Feature Map

# Max Pooling




Reduces the dimensionality  
without losing too much  
information

Reduces the number of  
parameters to train


Makes the network invariant to  
small transformations,  
distortions and translations




# Max Pooling



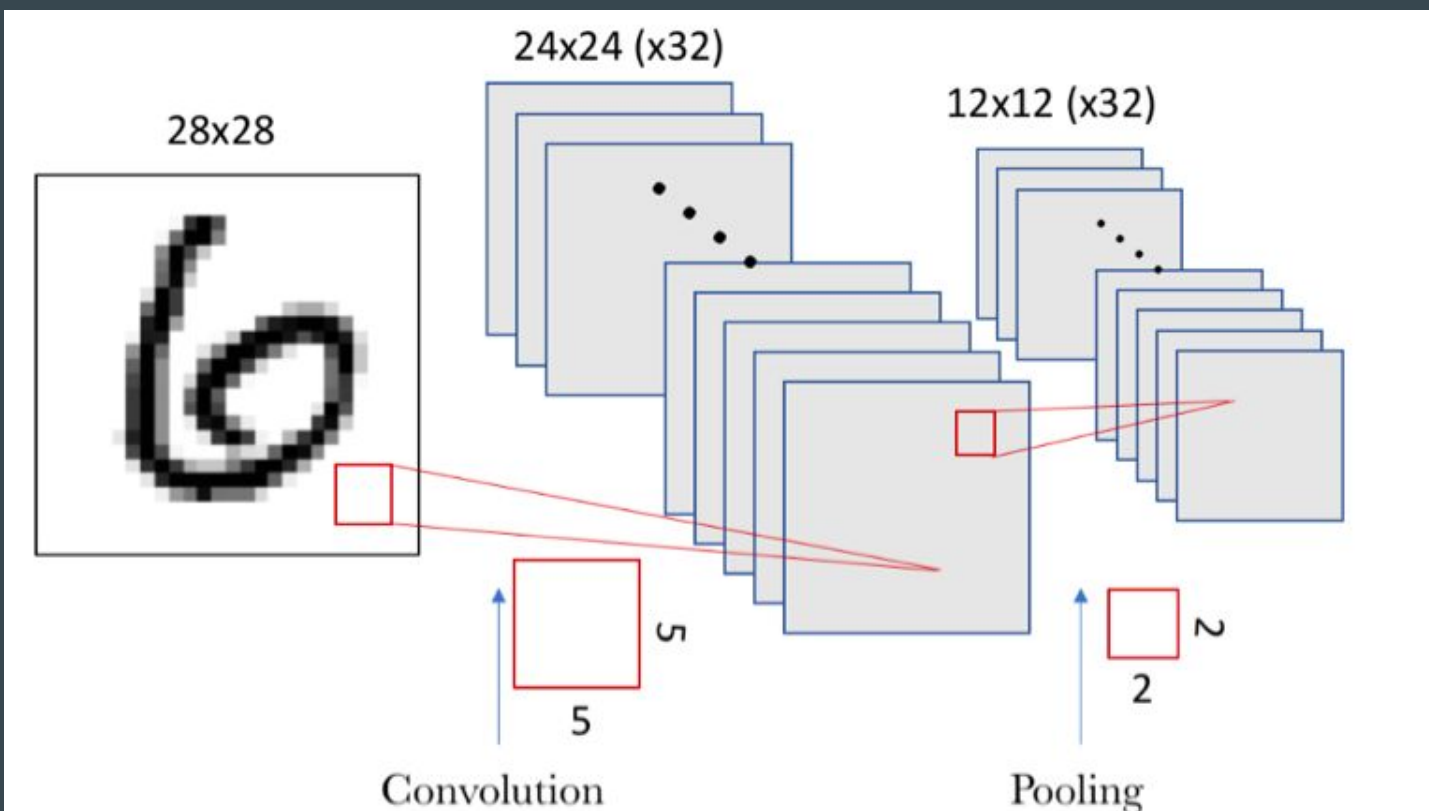
# Padding




Stride 1 with Padding



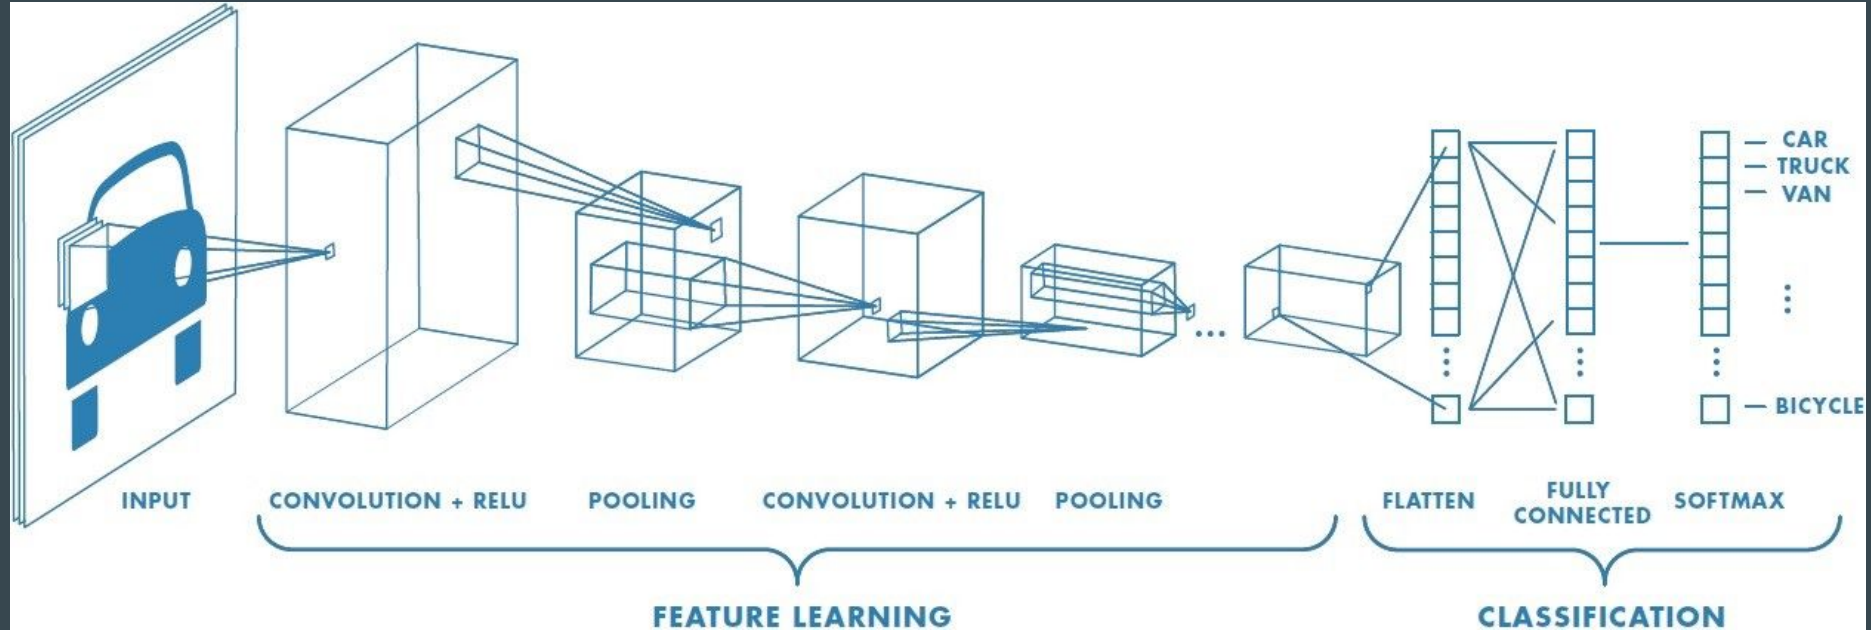
Feature Map




# Fully connected layer



# Piecing it all together





# Data Augmentation



Artificially increases data size

Prevents overfitting




# Dropout


Used for regularisation

Prevents overfitting

Can be applied to input and hidden layers - not final layer



No Dropout



With Dropout




# Example

[Online Demo](#)

# Object Detection



Test Image



Prediction from  
our network

No idea!?!

# You Only Look Once - YOLO




# Examples: Transfer Learning - object detection



# Examples: Image segmentation



# Examples: Counting craters to estimate the age of planetary surfaces



**Research Team:** A/Prof. Gretchen Benedix, A/Prof. Jonathan Paxman, Dr. Martin Towner, Dr. Anthony Lagain, Mr. Chris Norman, Prof. Tele Tan, Prof. Phil Bland

**CIC Specialists:** Dr. Kevin Chai, Shiv Meka




# Examples Rooftop segmentation to estimate urban sprawl



**Research Team:** Peiyu Li, Kexiang Xu, Dr. Mohammad Swapan, Dr. Cecilia Xia

**CIC Specialist:** Shiv Meka

# Examples: Benthic Habitat Mapping





# Detecting faults in Fin Fans from acoustic sensors



**Team:** Dr. Kristofer McKee, Dr. Amir Amin, Prof. Ian Howard, Jack Wiltshire (Cisco), Nathan Jombwe (Cisco)  
**CIC specialist:** Shiv Meka

# Web links

- <https://towardsdatascience.com/convolutional-neural-networks-for-all-part-i-cdd282ee7947>
- <https://medium.com/@ageitgey/machine-learning-is-fun-part-3-deep-learning-and-convolutional-neural-networks-f40359318721>
- <https://towardsdatascience.com/applied-deep-learning-part-4-convolutional-neural-networks-584bc134c1e2>
- <http://neuralnetworksanddeeplearning.com/chap5.html>
- <http://mathworld.wolfram.com/Convolution.html>
- <https://distill.pub/2017/feature-visualization/>